

Pandas Template

In [1]: `import pandas as pd`

In [2]: `# series: 1-dimentional use []#
series = pd.Series(["Cat", "Dog", "Donkey"])
hehavior=pd.Series(["Meow", "Bark", "Squeek"])
color=pd.Series(["Black", "White", "Grey"])`

In [3]: `# dataframe: 2-dimentional use {}#
animal = pd.DataFrame({"Animal":series, "Behavior":hehavior})
animal.head()`

Out[3]:

	Animal	Behavior
0	Cat	Meow
1	Dog	Bark
2	Donkey	Squeek

Import Data

In [4]: `#read data from static file#
carsales=pd.read_csv("car-sales-Copy1.csv")
carsales.to_csv("car-sales.csv", index=False) # so the index column will not be a column in the exported file`

In [5]: `#read data from GitHub, make sure GitHub in "raw" status#
heart_disease = pd.read_csv("https://raw.githubusercontent.com/mrdbourk/e/zero-to-mastery-ml/master/data/heart-disease.csv")`

Describe Data

In [6]: `#if a function, use (); if an attribute, don't use ()#
carsales.dtypes`

Out[6]:

Make	object
Colour	object
Odometer (KM)	int64
Doors	int64
Price	object
dtype:	object

```
In [7]: carsales.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Make                   10 non-null    object
1   Colour                 10 non-null    object
2   Odometer (KM)          10 non-null    int64
3   Doors                  10 non-null    int64
4   Price                  10 non-null    object
dtypes: int64(2), object(3)
memory usage: 528.0+ bytes
```

```
In [8]: #will show mean of numeric columns#
carsales.mean()
```

```
Out[8]: Odometer (KM)    78601.4
Doors                4.0
dtype: float64
```

```
In [9]: #will sum all columns#
carsales.sum()
```

```
Out[9]: Make           ToyotaHondaToyotaBMWNIssanToyotaHondaHondaToyo...
Colour                WhiteRedBlueBlackWhiteGreenBlueBlueWhiteWhite
Odometer (KM)                                786014
Doors                                           40
Price           $4,000.00$5,000.00$7,000.00$22,000.00$3,500.00...
dtype: object
```

```
In [10]: #query just one column entire column date using []#
carsales["Odometer (KM)"].sum()
```

```
Out[10]: 786014
```

```
In [11]: len(carsales)
```

```
Out[11]: 10
```

View and Select Data

```
In [12]: #show top and bottom 5 only#  
carsales.head()  
carsales.tail()
```

Out[12]:

	Make	Colour	Odometer (KM)	Doors	Price
5	Toyota	Green	99213	4	\$4,500.00
6	Honda	Blue	45698	4	\$7,500.00
7	Honda	Blue	54738	4	\$7,000.00
8	Toyota	White	60000	4	\$6,250.00
9	Nissan	White	31600	4	\$9,700.00

```
In [13]: #.loc & .iloc  
animal2 = pd.Series(["cat", "dog", "panda"], index=[0, 3, 5])  
animal2
```

Out[13]: 0 cat
3 dog
5 panda
dtype: object

```
In [14]: #.loc refers to index#  
animal2.loc[5]
```

Out[14]: 'panda'

```
In [15]: #.iloc refers to position 0,1,2#  
animal2.iloc[2]
```

Out[15]: 'panda'

```
In [16]: #.iloc slicing giving first 2 columns#  
animal2.iloc[:2]
```

Out[16]: 0 cat
3 dog
dtype: object

```
In [17]: #.loc slicing giving index up to 2#  
animal2.loc[:2]
```

Out[17]: 0 cat
dtype: object

```
In [18]: #select a column: two ways#
carsales["Make"]
carsales.Make
```

```
Out[18]: 0    Toyota
1     Honda
2    Toyota
3      BMW
4     Nissan
5    Toyota
6     Honda
7     Honda
8    Toyota
9     Nissan
Name: Make, dtype: object
```

```
In [19]: #filter with boolean criteria#
carsales[carsales["Make"]=="Toyota"]
carsales[carsales["Odometer (KM)"]>100000]
```

```
Out[19]:
```

	Make	Colour	Odometer (KM)	Doors	Price
0	Toyota	White	150043	4	\$4,000.00
4	Nissan	White	213095	4	\$3,500.00

```
In [20]: #!!! easy way to make a 2-dimentional pivot!!!#
pd.crosstab(carsales["Make"],carsales["Doors"])
```

```
Out[20]:
```

	Doors	3	4	5
Make				
BMW	0	0	1	
Honda	0	3	0	
Nissan	0	2	0	
Toyota	1	3	0	

```
In [21]: #groupby function, returns mean of all the numerical fields#
carsales.groupby(["Make"]).mean()
```

```
Out[21]:
```


	Odometer (KM)	Doors
Make		
BMW	11179.000000	5.00
Honda	62778.333333	4.00
Nissan	122347.500000	4.00
Toyota	85451.250000	3.75

Visualization

```
In [22]: #pip install matplotlib#
import matplotlib
import matplotlib.pyplot as plt
```


```
In [23]: carsales["Odometer (KM)"].plot()
```

```
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x11a370350>
```



```
In [24]: #run distribution#
carsales["Odometer (KM)"].hist()
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x11a6fc9d0>
```



```
In [25]: # Price is a string column #
# Convert string to integer#
# Use str.replace for a regex, lost the cents#
carsales["Price"]=carsales["Price"].str.replace('[\$,\.]', '').astype(int)
```

```
In [26]: carsales["Price"].dtype
```

```
Out[26]: dtype('int64')
```

Manipulating Data

```
In [27]: # This won't save to carsales#  
carsales["Make"].str.lower()
```

```
Out[27]: 0    toyota  
1    honda  
2    toyota  
3    bmw  
4    nissan  
5    toyota  
6    honda  
7    honda  
8    toyota  
9    nissan  
Name: Make, dtype: object
```

```
In [28]: # Save the format change  
carsales["Make"] = carsales["Make"].str.lower()  
carsales.head()
```

```
Out[28]:
```

	Make	Colour	Odometer (KM)	Doors	Price
0	toyota	White	150043	4	400000
1	honda	Red	87899	4	500000
2	toyota	Blue	32549	3	700000
3	bmw	Black	11179	5	2200000
4	nissan	White	213095	4	350000

```
In [29]: carmissingdata = pd.read_csv("car-sales-missing-data.csv")
carmissingdata
```

Out[29]:

	Make	Colour	Odometer	Doors	Price
0	Toyota	White	150043.0	4.0	\$4,000
1	Honda	Red	87899.0	4.0	\$5,000
2	Toyota	Blue	NaN	3.0	\$7,000
3	BMW	Black	11179.0	5.0	\$22,000
4	Nissan	White	213095.0	4.0	\$3,500
5	Toyota	Green	NaN	4.0	\$4,500
6	Honda	NaN	NaN	4.0	\$7,500
7	Honda	Blue	NaN	4.0	NaN
8	Toyota	White	60000.0	NaN	NaN
9	NaN	White	31600.0	4.0	\$9,700

```
In [30]: #fill N/A value with mean#
carmissingdata["Odometer"].fillna(carmissingdata["Odometer"].mean())
carmissingdata# see no change, have to use reassignment
```

Out[30]:

	Make	Colour	Odometer	Doors	Price
0	Toyota	White	150043.0	4.0	\$4,000
1	Honda	Red	87899.0	4.0	\$5,000
2	Toyota	Blue	NaN	3.0	\$7,000
3	BMW	Black	11179.0	5.0	\$22,000
4	Nissan	White	213095.0	4.0	\$3,500
5	Toyota	Green	NaN	4.0	\$4,500
6	Honda	NaN	NaN	4.0	\$7,500
7	Honda	Blue	NaN	4.0	NaN
8	Toyota	White	60000.0	NaN	NaN
9	NaN	White	31600.0	4.0	\$9,700

```
In [31]: #use reassignment
carmissingdata["Odometer"]=carmissingdata["Odometer"].fillna(carmissingdata["Odometer"].mean())
```

```
In [32]: # use inplace=True to avoid reassignment#
carmissingdata["Odometer"].fillna(carmissingdata["Odometer"].mean(),inplace=True)
```

```
In [33]: # use dropna to remove the rows that has NaN#
carmissingdata.dropna()
```

Out[33]:

	Make	Colour	Odometer	Doors	Price
0	Toyota	White	150043.000000	4.0	\$4,000
1	Honda	Red	87899.000000	4.0	\$5,000
2	Toyota	Blue	92302.666667	3.0	\$7,000
3	BMW	Black	11179.000000	5.0	\$22,000
4	Nissan	White	213095.000000	4.0	\$3,500
5	Toyota	Green	92302.666667	4.0	\$4,500

Creating Data

```
In [34]: #Add a Column from Series (not list)
#always append at the end
seats_column=pd.Series([5,5,5,5,5])
carsales["Seats"] = seats_column
carsales
```

Out[34]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats
0	toyota	White	150043	4	400000	5.0
1	honda	Red	87899	4	500000	5.0
2	toyota	Blue	32549	3	700000	5.0
3	bmw	Black	11179	5	2200000	5.0
4	nissan	White	213095	4	350000	5.0
5	toyota	Green	99213	4	450000	NaN
6	honda	Blue	45698	4	750000	NaN
7	honda	Blue	54738	4	700000	NaN
8	toyota	White	60000	4	625000	NaN
9	nissan	White	31600	4	970000	NaN


```
In [35]: #fillna with an exact value#
carsales["Seats"].fillna(5,inplace=True)
carsales
```

Out[35]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats
0	toyota	White	150043	4	400000	5.0
1	honda	Red	87899	4	500000	5.0
2	toyota	Blue	32549	3	700000	5.0
3	bmw	Black	11179	5	2200000	5.0
4	nissan	White	213095	4	350000	5.0
5	toyota	Green	99213	4	450000	5.0
6	honda	Blue	45698	4	750000	5.0
7	honda	Blue	54738	4	700000	5.0
8	toyota	White	60000	4	625000	5.0
9	nissan	White	31600	4	970000	5.0

```
In [36]: # Add a Column from List (not Series), need exact same amount of element
s in the new list#
fuel_economy=[3.2,2.3,4.4,5.0,6.0,2.4,7.7,8.1,6.3,5.4] # need to fill al
l the spaces, otherwise, error
carsales["fuel_economy"] = fuel_economy
carsales
```

Out[36]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats	fuel_economy
0	toyota	White	150043	4	400000	5.0	3.2
1	honda	Red	87899	4	500000	5.0	2.3
2	toyota	Blue	32549	3	700000	5.0	4.4
3	bmw	Black	11179	5	2200000	5.0	5.0
4	nissan	White	213095	4	350000	5.0	6.0
5	toyota	Green	99213	4	450000	5.0	2.4
6	honda	Blue	45698	4	750000	5.0	7.7
7	honda	Blue	54738	4	700000	5.0	8.1
8	toyota	White	60000	4	625000	5.0	6.3
9	nissan	White	31600	4	970000	5.0	5.4

```
In [37]: # Create a column with other columns' operation#
carsales["total_fuel_used"]=carsales["Odometer (KM)"]/100*carsales["fuel_economy"]
carsales.head()
```

Out[37]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats	fuel_economy	total_fuel_used
0	toyota	White	150043	4	400000	5.0	3.2	4801.376
1	honda	Red	87899	4	500000	5.0	2.3	2021.677
2	toyota	Blue	32549	3	700000	5.0	4.4	1432.156
3	bmw	Black	11179	5	2200000	5.0	5.0	558.950
4	nissan	White	213095	4	350000	5.0	6.0	12785.700

```
In [38]: # Create a column with one single value
carsales["wheels"]=4 # can put "True" any type in here
carsales.head()
```

Out[38]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats	fuel_economy	total_fuel_used	wheels
0	toyota	White	150043	4	400000	5.0	3.2	4801.376	4
1	honda	Red	87899	4	500000	5.0	2.3	2021.677	4
2	toyota	Blue	32549	3	700000	5.0	4.4	1432.156	4
3	bmw	Black	11179	5	2200000	5.0	5.0	558.950	4
4	nissan	White	213095	4	350000	5.0	6.0	12785.700	4

```
In [39]: # Remove the column, if you are talking about a column, axis=1
carsales=carsales.drop("wheels",axis=1)
carsales.head()
```

Out[39]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats	fuel_economy	total_fuel_used
0	toyota	White	150043	4	400000	5.0	3.2	4801.376
1	honda	Red	87899	4	500000	5.0	2.3	2021.677
2	toyota	Blue	32549	3	700000	5.0	4.4	1432.156
3	bmw	Black	11179	5	2200000	5.0	5.0	558.950
4	nissan	White	213095	4	350000	5.0	6.0	12785.700

Select Sample

```
In [40]: # shuffle and select sample =50% of the data if 0.5, 100% of the data if
frac=1
carsales.sample(frac=0.5)
carsales_shuffle=carsales.sample(frac=1)
carsales_shuffle.head()
```

Out[40]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats	fuel_economy	total_fuel_used
1	honda	Red	87899	4	500000	5.0	2.3	2021.677
2	toyota	Blue	32549	3	700000	5.0	4.4	1432.156
3	bmw	Black	11179	5	2200000	5.0	5.0	558.950
8	toyota	White	60000	4	625000	5.0	6.3	3780.000
6	honda	Blue	45698	4	750000	5.0	7.7	3518.746

```
In [41]: # put it back to order "reset_index"
# that this, there would be an index column, to remove the index column,
use drop=True
carsales_shuffle.reset_index(drop=True,inplace=True)
carsales_shuffle.head()
```

Out[41]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats	fuel_economy	total_fuel_used
0	honda	Red	87899	4	500000	5.0	2.3	2021.677
1	toyota	Blue	32549	3	700000	5.0	4.4	1432.156
2	bmw	Black	11179	5	2200000	5.0	5.0	558.950
3	toyota	White	60000	4	625000	5.0	6.3	3780.000
4	honda	Blue	45698	4	750000	5.0	7.7	3518.746

```
In [42]: #change a column with apply and lambda
carsales["Odometer (KM)"]=carsales["Odometer (KM)"].apply(lambda x: x/1.6)
carsales.head()
```

Out[42]:

	Make	Colour	Odometer (KM)	Doors	Price	Seats	fuel_economy	total_fuel_used
0	toyota	White	93776.875	4	400000	5.0	3.2	4801.376
1	honda	Red	54936.875	4	500000	5.0	2.3	2021.677
2	toyota	Blue	20343.125	3	700000	5.0	4.4	1432.156
3	bmw	Black	6986.875	5	2200000	5.0	5.0	558.950
4	nissan	White	133184.375	4	350000	5.0	6.0	12785.700

##